

Homework 1 – A scalable environment

Lab Information

Due Date:

Homework 1 Dropbox Deadline

Coordination:

This homework is to be completed alone

Objectives/Goal:

To understand popular incarnations of interception technologies that deal with web applications.

You have learned about various web technologies this week. Your goal is to set up an environment that is designed to scale with a large amount of load using this knowledge and knowledge from previous courses. As you design this, consider how each of these components interacts with HTTP Requests. Your goal will be to set up an environment with 3 different types of components: A Load Balancer, A caching server, and at least two Web Servers.

Deliverables:

- One zipped (ZIP only) file that contains 4 directories each with needed Dockerfile/compose.yaml (and associated files). See the following example of the layout. If you are unsure reach out to your instructor.
 - hw1-lastname.zip
 - act1/
 - Dockerfile
 - Docker-compse.yaml (optional on act 1)
 - act2/
 - cache/Dockerfile
 - webserver/Dockerfile
 - docker-compose.yaml
 - act3/
 - loadbalancer/Dockerfile
 - cache/Dockerfile
 - webserver/Dockerfile
 - docker-compose.yaml
 - act4/
 - step1/act4ste1.py
 - step1/requirements.txt
 - step2/act4ste2.py
 - step2/requirements.txt
 - Writeup.pdf
 - bonus-docker-compose.yml (optional)
 - A document (PDF/Word) with a breakdown in support of your choices.

Table of Contents:

Lab Information 1

Activity 1: Webserver Setup: 3

Activity 2: Caching Server Setup: 3

Activity 3: Load Balancer Setup..... 4

Activity 4: Proxy Scanning..... 5

Writeup..... 6

Deliverables 7

Activity 1: Webserver Setup:

Your first objective, if you choose to accept it, is to set up a web server. While Apache has the highest market share there are a number of possible, mature, alternatives. Webservers are simply servers that will respond to HTTP requests, although as we have gone over in class there are other supported protocols in many cases.

Step 1: Choose a webserver

Research available webservers and choose one which you think best suits the need of the remainder of the lab. Generally, you are given the choice of the following:

- Apache
- Nginx
- IIS

You may make, or choose a different technology with professor approval, however, the aforementioned technologies may be used without additional approval.

Step 2: setup a webserver

Now that you've made a decision on what platform you will use, it is time for you to install that webserver and configure it to serve a basic page. The configuration will vary based on which platform and OS you choose. Once you have the web server installed setup a basic text-based page that simply displays "Hello World" as part of the body tag. Your page's source must be verified as HTML 5 compliant using the W3C validator (<https://validator.w3.org>). This will be checked in your submitted Dockerfile. You may feel free to further customize your page as you see fit, ensuring you have met the following two requirements.

This webserver should be submitted as a Dockerfile (with optional compose). Docker is supported on Windows, MacOS, and Linux. If you are unfamiliar with Docker, you might check out this tutorial: <https://docker-curriculum.com/>

Activity 2: Caching Server Setup:

Caching servers are important staples of computing security. At each layer where speed changes occur, caches can be used to mitigate the performance losses if the right algorithm is used to take advantage of locality. Your second activity is to use such a technology to speed up your environment.

Step 1: Choose a caching server

Research available caching servers and choose one which you think best suits the need of the remainder of the lab. Generally, you are given the choice of the following:

- Varnish
- Nginx
- Squid

You may make, or choose a different technology with professor approval, however, the aforementioned technologies may be used without additional approval.

Step 2: setup a caching server

Now that you've made a decision on what platform you will use, you should install it. You should come up with a simple method to verify that the caching server is, in fact, working as this will be tested in your submitted work.

Hint: this may involve changing the page and demonstrating that the cache still displays on the web browser.

*Note: **The standard configuration for a web cache is as a reverse proxy.** This is the model you are expected to implement. This means that you will not have to change your local browser proxy settings*

You should create a new Docker image when creating your Caching Server. This Caching Server Docker image you've made should work in conjunction with the Web Server Docker image you made in the previous step. To do this you should use a Docker Compose. Docker Compose is a tool for setting up multi-container environments. on Windows, MacOS, and Linux. If you are unfamiliar with Docker Compose, you might check out this tutorial: <https://www.digitalocean.com/community/tutorials/how-to-install-and-use-docker-compose-on-ubuntu-20-04>. For an example see <https://github.com/dockersamples/example-voting-app>.

Activity 3: Load Balancer Setup

Now we are going to complicate things. Your task is going to be to do some research and determine the best location for a load balancer within the environment you already created. This will, in general, require you to add another webserver.

Note: Consulting your instructor is NOT considered research.

Step 1: Choose a load balancer

Research available caching servers and choose one which you think best suits the need of the remainder of the lab. Generally, you are given the choice of the following:

- Nginx
- HAProxy
- Varnish
- Apache
- Zen Load Balancer

You may make, or choose a different technology with professor approval, however, the aforementioned technologies may be used without additional approval.

Step 2: setup a load balancer

Now that you've chosen a load balancer, determine where you'll put it. Once you've figured out where it goes in your topology (web server and caching server so far), you might find you need another webserver – deploy that as needed. Finally, deploy your load balancer(s) and demonstrate it is functional by having it serve [slightly] different versions of the same page. Instead of saying 'Hello World', have the body of the other server say 'Hello World2'.

You should create a new Dockerfile when creating your load balancer. The load balancer Docker container you've made should work in conjunction with the web server and caching server Docker containers you made in the first and second step. You'll need at least two copies of the webserver.

Activity 4: Proxy Scanning

As we discussed in class HTTP has built-in proxy capabilities. Often these are capabilities are disabled, but if you can find a webserver that supports HTTP proxying, it can be very useful for surfing the web anonymously. In this exercise, you will be using an existing library to try and find these open proxies.

Step 1: Using the requests library (or equivalent).

If you've never used [Python's Requests library](#) now would be a great time to become familiar with it. The library makes it very easy to make standard HTTP requests to a site without having to know much about the underlying protocol. In the remainder of the course, we will be creating a similar library. For this step you'll simply need to leverage the Requests library to request <https://csec.rit.edu>.

Submit this step as part of your zip in a folder called 'step1' that contains a Python script named `act4ste1.py` and a pip-compliant `requirements.txt` that outlines which libraries your script requires.

Step 2: Write a scanner that will try to find an anonymous HTTP proxy

Now that you have made a basic HTTP Requests script, all you're going to need to do is slightly enhance it. You should write a script such that it can take in a range of IP's and scan each one looking for an open HTTP proxy.

We talked briefly about how these requests look but you will only need to know how to use Python Requests (or similar) in order to generate it (hint: no need for a raw HTTP packet). Once you have it correctly generating scan ranges of the internet until you find an open proxy server. Your script should return the IP addresses (one per line) of the Proxies discovered

Note: A proxy server is going to respond differently from a stock webserver in various ways. Your goal is to figure out how to determine it is actually proxying traffic.

Hint: You may want to setup a proxy server or try using a public list like <http://www.freeproxylists.net/> (please remember many of these are going to malicious).

Hint: You may wish to leverage threading to speed this up.

Submit this step as part of your zip in a folder called 'step2' that contains a Python script named act4ste2.py and a requirements.txt that outlines which libraries your script requires. Note: act4ste2.py should take two command line arguments a start IP and an end IP.

Writeup

As you go through the remainder of your career you will often be faced with choosing a single technology among many choices. In this case, you were given a set of options, all of which were free(ish) – but there are also many situations where the cost will factor in. Given the above options, how did you select which technologies you would use? Do some of these combinations work better together than others? Are some easier to use or deploy?

Also, attach your topology and discuss why you chose this topology and what the alternatives might have been. Lastly, include a discussion about why these technologies might be needed. At the same time, you may wish to discuss why the single text web page example we just created might not need all the complexity we developed.

Note: if you've typed more than a paragraph for each answer, you've written too much.

Submit a writeup in PDF or Microsoft Word format. This document should answer the above points and have an image of your topology. Include this writeup in your zip submission.

Bonus

Containerization is a powerful technology used by nearly all major organizations to ensure consistency and scalability of their environments. Initially, it may seem like containerizing content has a higher upfront cost, but much like object orientation, if reuse and portability is intended it can drastically speed up processes in the long run.

With that in mind, there are often less verbose ways to express configurations that are designed to enhance portability. In Docker this is often done by leveraging existing containers, from Docker registries, and by limiting the amount of additional associated files.

To get the bonus you must create a docker-compose.yml file that is able to describe the first three activities of this homework (2 Webservers, Load Balancer, and Caching server and the associated configurations) using ONLY a docker-compose.yaml (no other files or folders are allowed).

Submit a file bonus-docker-compose.yml with your zip submissions.

Deliverables

Activity 1 – Docker/Compose files for a functional webserver (20%)

Activity 2 – Docker/compose files for a Caching Server as part of Docker Compose environment (20%)

Activity 3 – Docker/compose files for a Load Balancer as part of Docker Compose environment (20%)

Activity 4 – Two different Python scripts – one from Step 1 and one from Step 2 (part 1 – 10%, part 2 – 20%)

Writeup (10%)

Bonus [Optional] (10%)